

ASDS 6304 - Project 2

Soham Neeraj Agarkar & Utkarsh Pant on September 27, 2024.

Part 1. PCA on the Wine Dataset

Task 1.1. Data Preparation

Load the Wine dataset and normalize the features. Why is feature normalization important when applying PCA?

```
In [1]: from sklearn.decomposition import PCA, KernelPCA, SparsePCA
        from sklearn.preprocessing import StandardScaler
        from sklearn.datasets import load_digits
        import pandas as pd
        import numpy as np

        import matplotlib.pyplot as plt

In [2]: wine_data = pd.read_csv('./wine_dataset-1.csv')
        wine_data.columns = wine_data.columns.str.replace(' ', '_').str.lower()
        x = wine_data.drop('class', axis=1)
        y = wine_data['class']
```

Normalization of features is important when applying PCA, since the difference in the scale of each feature may cause it to have an erroneously larger or smaller projection in lower dimensions, thus skewing the results incorrectly.

Question to consider: what is the difference between "normalization" and "standardization," and when should we prefer one over the other?

We're using the **Standard Scaler** to scale our data to approximately match the Standard Normal distribution.

```
In [3]: scaler = StandardScaler()
        scaler.fit(x)
        wine_data_scaled = scaler.transform(x)

        wine_data_scaled

Out[3]: array([[ 1.51861254, -0.5622498 ,  0.23205254, ...,  0.36217728,
                  1.84791957,  1.01300893],
               [ 0.24628963, -0.49941338, -0.82799632, ...,  0.40605066,
                  1.1134493 ,  0.96524152],
               [ 0.19687903,  0.02123125,  1.10933436, ...,  0.31830389,
                  0.78858745,  1.39514818],
               ...,
               [ 0.33275817,  1.74474449, -0.38935541, ..., -1.61212515,
                  -1.48544548,  0.28057537],
               [ 0.20923168,  0.22769377,  0.01273209, ..., -1.56825176,
                  -1.40069891,  0.29649784],
               [ 1.39508604,  1.58316512,  1.36520822, ..., -1.52437837,
                  -1.42894777, -0.59516041]])
```

Now that the data has been scaled, we can apply PCA to it.

Task 1.2. Apply PCA

Apply PCA to reduce the dimensionality of the Wine dataset. Keep enough principal components to explain 90% of the variance. How many principal components are required to capture 90% of the variance?

```
In [4]: pca = PCA(n_components=0.9)
        pca.fit(wine_data_scaled)
        wine_data_pca = pca.transform(wine_data_scaled)

        wine_data_pca.shape

Out[4]: (178, 8)
```

From the above result, we can see that we need **8** features to explain 90% of the variance in the data. This was achieved by setting `n_components=0.9` in the `PCA` constructor, which means that the solver automatically determines the right number of principal components to explain the given percentage of variance.

```
In [5]: # Kernel PCA on the wine dataset
kpca = KernelPCA(n_components=8, kernel='rbf', fit_inverse_transform=True)
wine_data_kpca = kpca.fit_transform(wine_data_scaled)

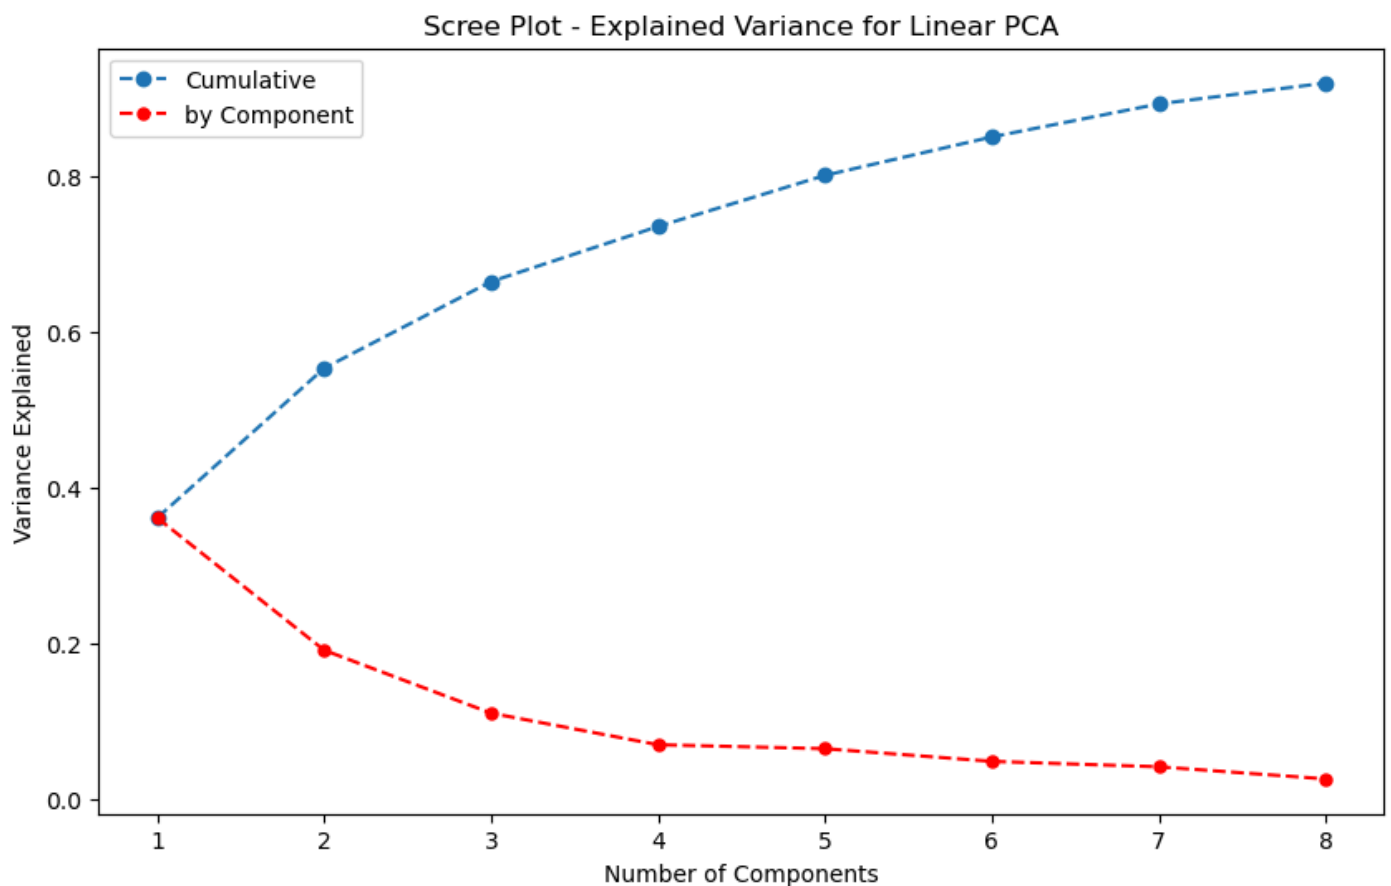
# Sparse PCA on the wine dataset
spca = SparsePCA(n_components=8)
wine_data_spca = spca.fit_transform(wine_data_scaled)
```

Plot the explained variance ratio and the cumulative explained variance. What does this tell you about the dataset?

```
In [6]: explained_variance_ratio = pca.explained_variance_ratio_
cumulative_evr = np.cumsum(explained_variance_ratio)

plt.style.use('default')
plt.figure(figsize=(10, 6))
plt.plot(range(1, len(cumulative_evr) + 1),
         cumulative_evr, marker='o', linestyle='--')
# adding scree plot for variance explained by each component
plt.plot(range(1, len(explained_variance_ratio) + 1), explained_variance_ratio,
         marker='o', markersize=5, color='red', linestyle='--')
plt.title('Scree Plot - Explained Variance for Linear PCA')
plt.xlabel('Number of Components')
plt.ylabel('Variance Explained')
plt.legend(['Cumulative', 'by Component'])

plt.show()
```



The above Scree Plot shows us that each principal component in the data captures slightly more variance in the data than the lower-dimensional component, and PCA is effective for the Wine Dataset, since we go from about **14** to **8** features - a 43% reduction in dimensionality!

Task 1.3. Visualize and Interpret Results


Plot the data in the first two principal component spaces. Can you see clear separation between the three wine cultivars?

From the Scree plot above, we can see that PC1 and PC2 explain the greatest variance between them, so we can use them to visualize our results. This plot should show the clearest class differentiation!

```
In [7]: reduced_wine_data = pd.DataFrame(wine_data_pca, columns=[f'PC{i}' for i in range(1, wine_data_pca.shape[1] + 1)])
reduced_wine_data['class'] = y

plt.figure(figsize=(10, 6))
for class_value in y.unique():
    subset = reduced_wine_data[reduced_wine_data['class'] == class_value]
    plt.scatter(x=subset['PC1'], y=subset['PC2'], label=class_value)
plt.legend()
plt.title('PC1 vs. PC2 for the Wine Dataset, After Linear PCA. Class differentiation is visible.')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

plt.show()
```



Reconstruct the original dataset from the reduced PCA components. What information is lost when reducing the dimensionality?

```
In [8]: reconstructed_wine_data_lzca = pca.inverse_transform(wine_data_pca)
reconstructed_wine_data_kzca = kpca.inverse_transform(wine_data_kzca)
reconstructed_wine_data_spca = spca.inverse_transform(wine_data_spca)

reconstruction_error_wine_lzca = np.mean(np.square(wine_data_scaled - reconstructed_wine_data_lzca))
reconstruction_error_wine_kzca = np.mean(np.square(wine_data_scaled - reconstructed_wine_data_kzca))
reconstruction_error_wine_spca = np.mean(np.square(wine_data_scaled - reconstructed_wine_data_spca))
print(f'Reconstruction Error (Mean Square Error) Linear PCA: {reconstruction_error_wine_lzca}')
print(f'Reconstruction Error (Mean Square Error) Kernel PCA: {reconstruction_error_wine_kzca}')
print(f'Reconstruction Error (Mean Square Error) Sparse PCA: {reconstruction_error_wine_spca}')
```

Reconstruction Error (Mean Square Error) Linear PCA: 0.0798245565422737
Reconstruction Error (Mean Square Error) Kernel PCA: 0.34051048090707653
Reconstruction Error (Mean Square Error) Sparse PCA: 0.09122242442595138


Part 2. Kernel PCA on the Digits Dataset

Task 2.1. Data Preparation

Load the Digits dataset using `sklearn.datasets.load_digits()`. Visualize some of the digits using `matplotlib`.

```
In [9]: digits = load_digits()
digits_data = digits.data
digits_target = digits.target
```

```
In [10]: # visualise the first 10 digits
fig, ax = plt.subplots(2, 5, figsize=(10, 4))
plt.style.use('default')
plt.suptitle('First 10 digits of the Digits Dataset (8x8 images)')
for i in range(10):
    plt.subplot(2, 5, i + 1)
    plt.imshow(digits_data[i].reshape(8, 8), cmap='gray')
    plt.title(digits_target[i])
    plt.axis('off')
plt.show()
```



Task 2.2. Apply Kernel PCA with the RBF Kernel.

Apply Kernel PCA using an RBF kernel to reduce the dimensionality of the dataset.

```
In [11]: # scale the data
scaler = StandardScaler()
scaler.fit(digits_data)
digits_data_scaled = scaler.transform(digits_data)

# apply Kernel PCA
kPCA = KernelPCA(kernel='rbf', gamma=0.01, fit_inverse_transform=True)
digits_kPCA = kPCA.fit_transform(digits_data_scaled)


# apply linear and sparse PCA
pca = PCA(n_components=0.9)
digits_lPCA = pca.fit_transform(digits_data_scaled)
spca = SparsePCA(n_components=None)
digits_spca = spca.fit_transform(digits_data_scaled)
```

Visualize the data in the 2D space of the first two kernel principal components. What do you observe?

```
In [12]: # visualizing the first 2 PC's in 2D space

plt.figure(figsize=(10, 6))
plt.scatter(x=digits_kPCA[:, 0], y=digits_kPCA[:, 1], c=digits_target, cmap='viridis')
plt.title('First 2 Principal Components of the Digits Dataset, After Kernel PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar()
plt.show()
```

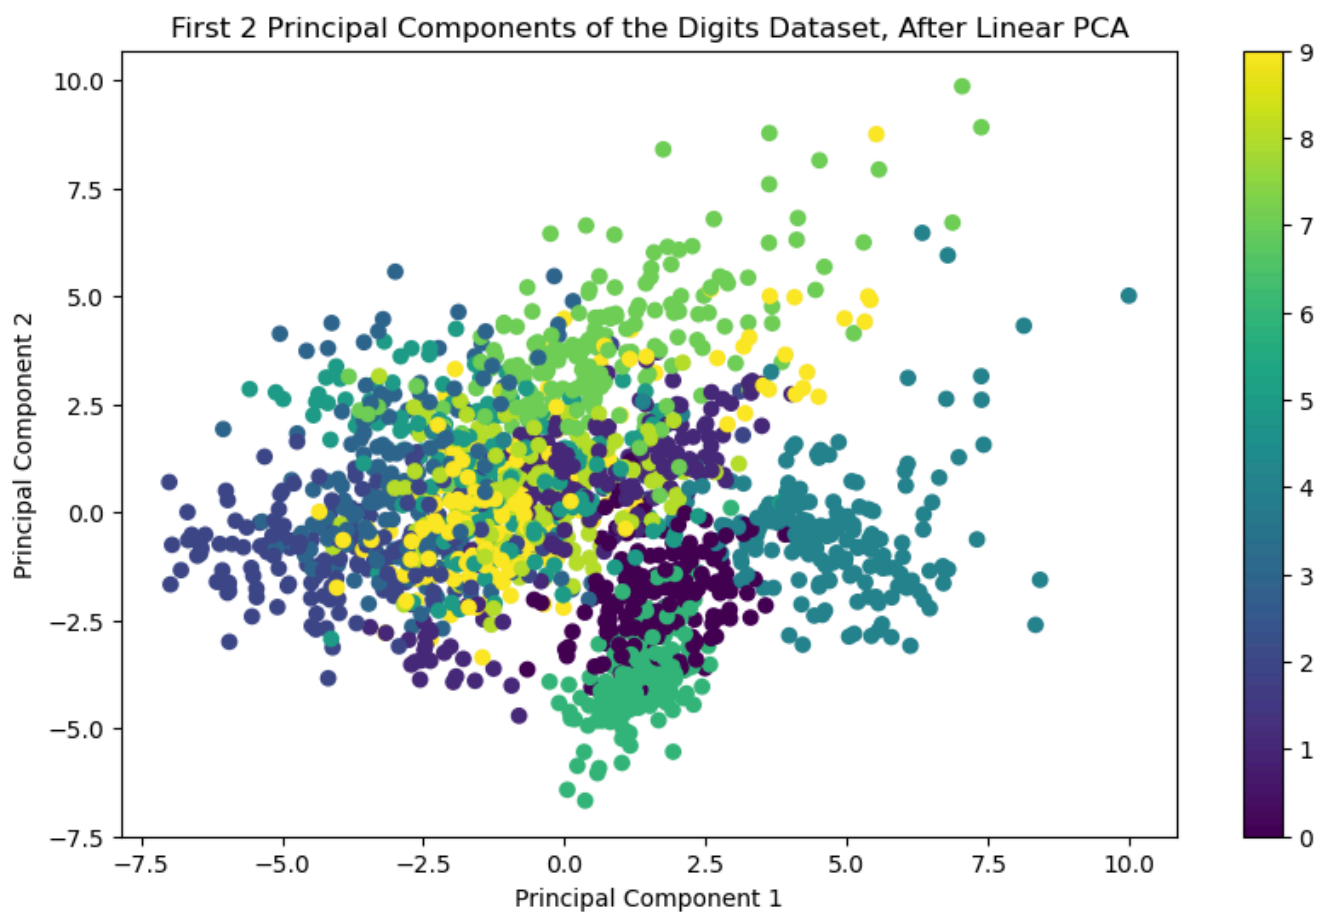
First 2 Principal Components of the Digits Dataset, After Kernel PCA



Task 2.3. Compare with Linear PCA

Apply standard PCA to reduce the data to 2 components. Plot the results and compare them with Kernel PCA results.

```
In [13]: plt.figure(figsize=(10, 6))
plt.scatter(x=digits_lpca[:, 0], y=digits_lpca[:, 1], c=digits_target, cmap='viridis')
plt.title('First 2 Principal Components of the Digits Dataset, After Linear PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar()
plt.show()
```



Here, we can see that Kernel PCA was able to obtain **much better class differentiation than linear PCA**.

Task 2.4. Interpretation and Evaluation.

Reconstruct the original data from the reduced components. Compare the reconstruction errors for both methods.

Referenced discussion

[Stats StackExchange: How to choose a kernel for kernel PCA?](#)

```
In [14]: digits_reconstructed_kpca = kpca.inverse_transform(digits_kpca)
digits_reconstructed_lpca = pca.inverse_transform(digits_lpca)
digits_reconstructed_spca = spca.inverse_transform(digits_spca)

reconstruction_error_digits_kpca = np.mean(np.square(digits_data_scaled - digits_reconstructed_kpca))
reconstruction_error_digits_lpca = np.mean(np.square(digits_data_scaled - digits_reconstructed_lpca))
reconstruction_error_digits_spca = np.mean(np.square(digits_data_scaled - digits_reconstructed_spca))

print(f'Reconstruction Error (Mean Square Error) for Kernel PCA: {reconstruction_error_digits_kpca}')
print(f'Reconstruction Error (Mean Square Error) for PCA: {reconstruction_error_digits_lpca}')
print(f'Reconstruction Error (Mean Square Error) for Sparse PCA: {reconstruction_error_digits_spca}')
```

Reconstruction Error (Mean Square Error) for Kernel PCA: 0.4721858946738644

Reconstruction Error (Mean Square Error) for PCA: 0.09487000241752752

Reconstruction Error (Mean Square Error) for Sparse PCA: 9.114453776024784e-05

Part 3: Sparse PCA on the Breast Cancer Dataset

Task 3.1. Data Preparation

Load and normalize the Breast Cancer dataset. Why might Sparse PCA be useful in this context?

```
In [15]: breast_cancer_data = pd.read_csv('./breast_cancer_dataset.csv')

breast_cancer_data.columns = breast_cancer_data.columns.str.replace(
    '_', '_').str.lower()
# drop the id and diagnosis columns for x
x = breast_cancer_data.drop(['id', 'diagnosis'], axis=1)
```

```
y = breast_cancer_data['diagnosis']
x.head()
```

```
Out[15]:
```

	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	feature_8	feature_9	feature_10	...	feature_20
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	25.3
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	24.9
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	23.5
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	14.9
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	22.5

5 rows × 30 columns



```
In [16]: # standard scale the data
scaler = StandardScaler()
scaler.fit(x)
breast_cancer_data_scaled = scaler.transform(x)
```

Applying Sparse PCA on the dataset after scaling it:

```
In [17]: sparse_pca = SparsePCA(n_components=2, alpha=0.1)
sparse_pca.fit(breast_cancer_data_scaled)
breast_cancer_data_sparse_pca = sparse_pca.transform(breast_cancer_data_scaled)
```

Applying Linear and Kernel PCA for comparison:

```
In [18]: linear_pca = PCA(n_components=2)
linear_pca.fit(breast_cancer_data_scaled)
breast_cancer_data_pca = linear_pca.transform(breast_cancer_data_scaled)

KernelPCA = KernelPCA(n_components=2, kernel='rbf', gamma=0.1)
KernelPCA.fit(breast_cancer_data_scaled)
breast_cancer_data_kpca = KernelPCA.transform(breast_cancer_data_scaled)
```

Task 3.4. Reconstruction and Comparison

Reconstruct the original dataset from the sparse principal components. Compare the reconstruction error with that of regular PCA.

```
In [19]: reconstructed_breast_cancer_data_spca = sparse_pca.inverse_transform(
    breast_cancer_data_sparse_pca)

reconstructed_breast_cancer_data_l pca = linear_pca.inverse_transform(
    breast_cancer_data_pca)

reconstructed_breast_cancer_data_kpca = linear_pca.inverse_transform(
    breast_cancer_data_kpca)

reconstruction_error_breast_cancer_sparse_pca = np.mean(
    np.square(breast_cancer_data_scaled - reconstructed_breast_cancer_data_spca))
print(f'Reconstruction Error (Mean Square Error) for Sparse PCA: {
    reconstruction_error_breast_cancer_sparse_pca}')

reconstruction_error_breast_cancer_l pca = np.mean(
    np.square(breast_cancer_data_scaled - reconstructed_breast_cancer_data_l pca))
print(f'Reconstruction Error (Mean Square Error) for PCA: {
    reconstruction_error_breast_cancer_l pca}')

reconstruction_error_breast_cancer_kpca = np.mean(
    np.square(breast_cancer_data_scaled - reconstructed_breast_cancer_data_kpca))
print(f'Reconstruction Error (Mean Square Error) for PCA: {
    reconstruction_error_breast_cancer_kpca}')
```

```
Reconstruction Error (Mean Square Error) for Sparse PCA: 0.36763944298781615
Reconstruction Error (Mean Square Error) for PCA: 0.3675679234844056
Reconstruction Error (Mean Square Error) for PCA: 1.0375589702179815
```

In this case, both methods achieve a similar error in reconstructed data.

Part 4: General Comparison of PCA, Kernel PCA, and Sparse PCA

Task 4.1. Summary Comparison

Compare the reconstruction errors for PCA, Kernel PCA, and Sparse PCA across the three datasets.

```
In [20]: # reconstruction error comparison table across all datasets
reconstruction_error_data = {
    'Dataset': ['Wine', 'Digits', 'Breast Cancer'],
    'PCA': [reconstruction_error_wine_lpca, reconstruction_error_digits_lpca, reconstruction_error_breast_cancer_lpca],
    'Kernel PCA': [reconstruction_error_wine_kpca, reconstruction_error_digits_kpca, reconstruction_error_breast_cancer_kpca],
    'Sparse PCA': [reconstruction_error_wine_spca, reconstruction_error_digits_spca, reconstruction_error_breast_cancer_spca]
}

reconstruction_errors = pd.DataFrame(reconstruction_error_data)
reconstruction_errors
```

```
Out[20]:
```

	Dataset	PCA	Kernel PCA	Sparse PCA
0	Wine	0.079825	0.340510	0.091222
1	Digits	0.094870	0.472186	0.000091
2	Breast Cancer	0.367568	1.037559	0.367639

Task 4.2. Visualization Comparison

For each dataset, visualize the first two principal components (or kernel components) obtained through PCA, Kernel PCA, and Sparse PCA.

```
In [32]: # wine dataset PC1 vs PC2 for linear, kernel and sparse PCA
# 1x3 grid

# Linear pca
fig, ax = plt.subplots(3, 3, figsize=(15, 17))
plt.subplot(3, 3, 1)
plt.scatter(x=reduced_wine_data['PC1'], y=reduced_wine_data['PC2'], c=reduced_wine_data['class'], cmap='viridis')
plt.title('Wine Dataset - Linear PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# wine kernel pca
plt.subplot(3, 3, 2)
plt.scatter(x=wine_data_kpca[:, 0], y=wine_data_kpca[:, 1], c=reduced_wine_data['class'], cmap='viridis')
plt.title('Wine Dataset - Kernel PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# wine sparse pca
plt.subplot(3, 3, 3)
plt.scatter(x=wine_data_spca[:, 0], y=wine_data_spca[:, 1], c=reduced_wine_data['class'], cmap='viridis')
plt.title('Wine Dataset - Sparse PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# digits linear pca
plt.subplot(3, 3, 4)
plt.scatter(x=digits_lpca[:, 0], y=digits_lpca[:, 1], c=digits_target, cmap='viridis')
plt.title('Digits Dataset - Linear PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# digits kernel pca
plt.subplot(3, 3, 5)
plt.scatter(x=digits_kpca[:, 0], y=digits_kpca[:, 1], c=digits_target, cmap='viridis')
plt.title('Digits Dataset - Kernel PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# digits sparse pca
plt.subplot(3, 3, 6)
# Convert 'diagnosis' to numerical values
```



```

diagnosis_numeric = breast_cancer_data['diagnosis'].map({'M': 1, 'B': 0})

plt.scatter(x=digits_spca[:, 0], y=digits_spca[:, 1], c=digits_target, cmap='viridis')
plt.title('Digits Dataset - Sparse PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

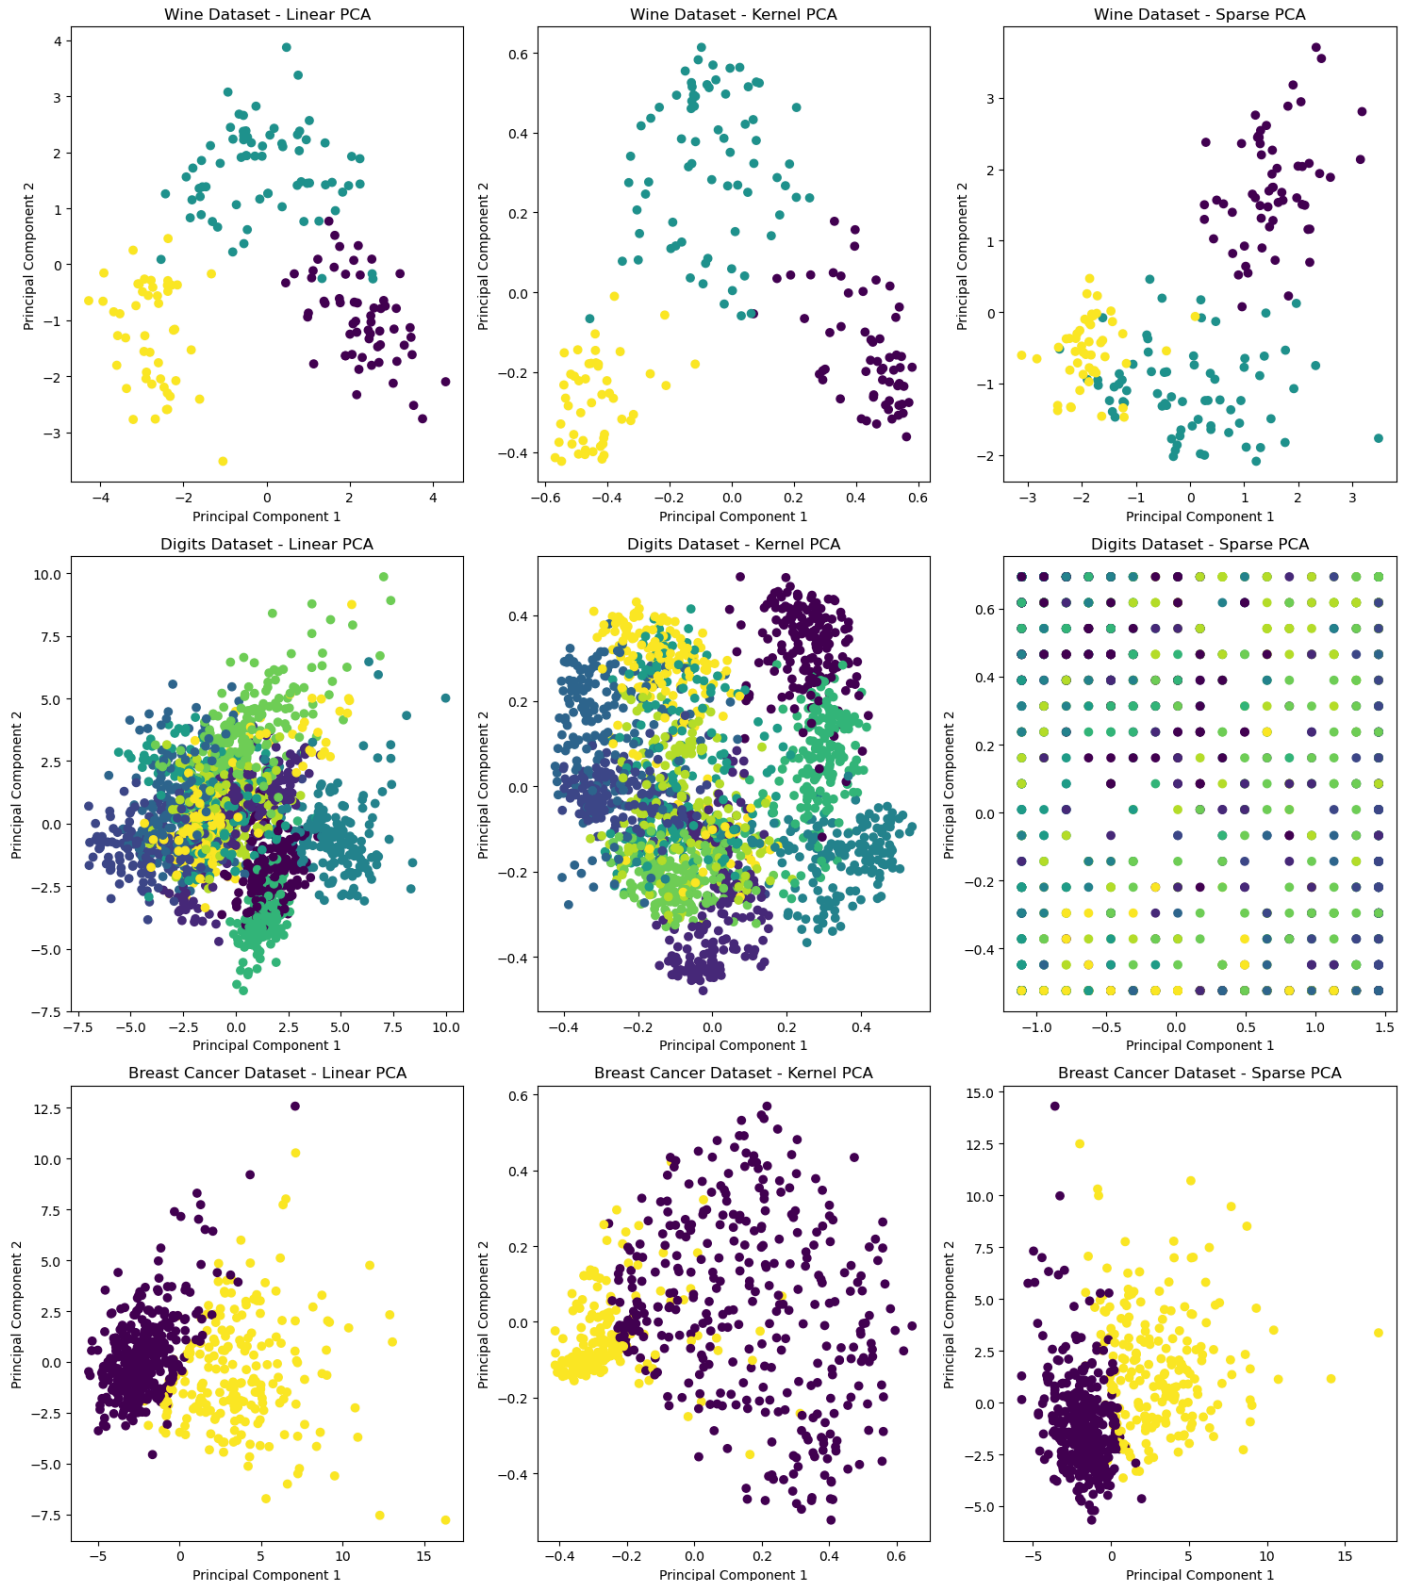
# breast cancer linear pca
plt.subplot(3, 3, 7)
plt.scatter(x=breast_cancer_data_pca[:, 0], y=breast_cancer_data_pca[:, 1], c=diagnosis_numeric, cmap='viridis')
plt.title('Breast Cancer Dataset - Linear PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# breast cancer kernel pca
plt.subplot(3, 3, 8)
plt.scatter(x=breast_cancer_data_kpca[:, 0], y=breast_cancer_data_kpca[:, 1], c=diagnosis_numeric, cmap='viridis')
plt.title('Breast Cancer Dataset - Kernel PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

# breast cancer sparse pca
plt.subplot(3, 3, 9)
plt.scatter(x=breast_cancer_data_sparse_pca[:, 0], y=breast_cancer_data_sparse_pca[:, 1], c=diagnosis_numeric, cm
plt.title('Breast Cancer Dataset - Sparse PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

plt.tight_layout()
plt.show()

```



Part 5: Principal Component Regression (PCR) on Boston Housing Dataset

Task 5.1. Data Preparation

Load the Boston Housing dataset, normalize the features, and apply PCA. How many components are needed to explain 95% of the variance in the dataset?

```
In [33]: boston_data = pd.read_csv('./boston_housing_dataset.csv')

boston_data.columns = boston_data.columns.str.replace(' ', '_').str.lower()

x = boston_data.drop('medv', axis=1)
y = boston_data['medv']

scaler = StandardScaler()
scaler.fit(x)
boston_data_scaled = scaler.transform(x)
```

```
pca = PCA(n_components=0.95)
pca.fit(boston_data_scaled)
boston_data_pca = pca.transform(boston_data_scaled)
```

```
In [34]: boston_data_pca.shape
```

```
Out[34]: (506, 9)
```

Thus, **9** Principal Components are needed to explain 95% of the variance in the dataset.

Task 5.2. Apply PCR

Use the principal components obtained from PCA to fit a linear regression model. Compare the performance of PCR with a regular linear regression model using the same dataset. Report the Mean Squared Error (MSE) and R^2 score for both models.

```
In [37]: # fit a principal component regression model
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

pca_regression_model = LinearRegression()
pca_regression_model.fit(boston_data_pca, y)

# calculate the R^2 score and RMSE
y_pred_pca = pca_regression_model.predict(boston_data_pca)
r2_pca = r2_score(y, y_pred_pca)
rmse_pca = np.sqrt(mean_squared_error(y, y_pred_pca))

print(f'R^2 Score: {r2_pca}', f'RMSE: {rmse_pca}', sep='\n')
```

```
R^2 Score: 0.7056708975511301
```

```
RMSE: 4.984689778970822
```

```
In [38]: # regular linear regression
linear_regression = LinearRegression()
linear_regression.fit(x, y)

y_pred = linear_regression.predict(x)
r2 = r2_score(y, y_pred)
rmse = np.sqrt(mean_squared_error(y, y_pred))

print(f'R^2 Score: {r2}', f'RMSE: {rmse}', sep='\n')
```

```
R^2 Score: 0.7406426641094095
```

```
RMSE: 4.679191295697281
```

In this case, it appears that the linear regression model using the features as-is, performs better than the PCR model.

This is confirmed by the fact that both, the R^2 score and the RMSE are better for the linear regression model - which means that it fits the data better. However, a better approach would be to split the datasets into training and testing data, and conduct validation of each model on the test datasets.